

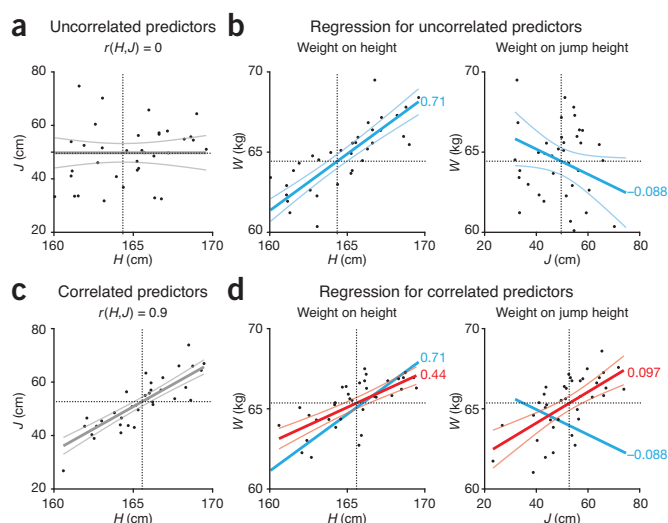
## POINTS OF SIGNIFICANCE

## Multiple linear regression

When multiple variables are associated with a response, the interpretation of a prediction equation is seldom simple.

Last month we explored how to model a simple relationship between two variables, such as the dependence of weight on height<sup>1</sup>. In the more realistic scenario of dependence on several variables, we can use multiple linear regression (MLR). Although MLR is similar to linear regression, the interpretation of MLR correlation coefficients is confounded by the way in which the predictor variables relate to one another.

In simple linear regression<sup>1</sup>, we model how the mean of variable  $Y$  depends linearly on the value of a predictor variable  $X$ ; this relationship is expressed as the conditional expectation  $E(Y|X) = \beta_0 + \beta_1 X$ . For more than one predictor variable  $X_1, \dots, X_p$ , this becomes  $\beta_0 + \sum \beta_j X_j$ . As for simple linear regression, one can use the least-squares estimator (LSE) to determine estimates  $b_j$  of the  $\beta_j$  regression parameters by minimizing the residual sum of squares,  $SSE = \sum (y_i - \hat{y}_i)^2$ , where  $\hat{y}_i = b_0 + \sum_j b_j x_{ij}$ . When we use the regression sum of squares,  $SSR = \sum (\hat{y}_i - \bar{y})^2$ , the ratio  $R^2 = SSR / (SSR + SSE)$  is the amount of variation explained by the regression model and in multiple regression is called the coefficient of determination.



**Figure 1** | The results of multiple linear regression depend on the correlation of the predictors, as measured here by the Pearson correlation coefficient  $r$  (ref. 2). (a) Simulated values of uncorrelated predictors,  $r(H,J) = 0$ . The thick gray line is the regression line, and thin gray lines show the 95% confidence interval of the fit. (b) Regression of weight ( $W$ ) on height ( $H$ ) and of weight on jump height ( $J$ ) for uncorrelated predictors shown in a. Regression slopes are shown ( $b_H = 0.71$ ,  $b_J = -0.088$ ). (c) Simulated values of correlated predictors,  $r(H,J) = 0.9$ . Regression and 95% confidence interval are denoted as in a. (d) Regression (red lines) using correlated predictors shown in c. Light red lines denote the 95% confidence interval. Notice that  $b_J = 0.097$  is now positive. The regression line from b is shown in blue. In all graphs, horizontal and vertical dotted lines show average values.

**Table 1** | Regression coefficients and  $R^2$  for different predictors and predictor correlations

Predictors in model	Regression coefficients			
	$\beta_H$	$\beta_J$	$\beta_0$	
$H, J$	0.7	-0.08	-46.5	
Predictors fitted	Estimated regression coefficients			
	$b_H$	$b_J$	$b_0$	$R^2$
Uncorrelated predictors, $r(H,J) = 0$				
$H$	0.71		-51.7	0.66
$J$		-0.088	69.3	0.19
$H, J$	0.71	-0.088	-47.3	0.85
Correlated predictors, $r(H,J) = 0.9$				
$H$	0.44		-8.1 (ns)	0.64
$J$		0.097	60.2	0.42
$H, J$	0.63	-0.056	-36.2	0.67

Actual ( $\beta_H, \beta_J, \beta_0$ ) and estimated regression coefficients ( $b_H, b_J, b_0$ ) and coefficient of determination ( $R^2$ ) for uncorrelated and highly correlated predictors in scenarios where either  $H$  or  $J$  or both  $H$  and  $J$  predictors are fitted in the regression. Regression coefficient estimates for all values of predictor sample correlation,  $r(H,J)$  are shown in Figure 2. ns, not significant.

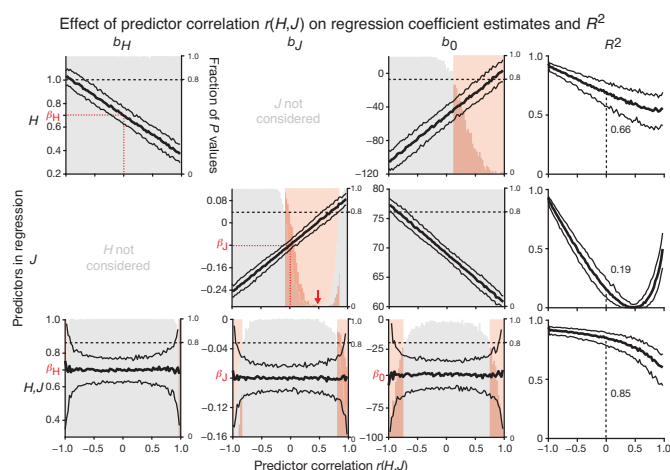
The slope  $\beta_j$  is the change in  $Y$  if predictor  $j$  is changed by one unit and others are held constant. When normality and independence assumptions are fulfilled, we can test whether any (or all) of the slopes are zero using a  $t$ -test (or regression  $F$ -test). Although the interpretation of  $\beta_j$  seems to be identical to its interpretation in the simple linear regression model, the innocuous phrase “and others are held constant” turns out to have profound implications.

To illustrate MLR—and some of its perils—here we simulate predicting the weight ( $W$ , in kilograms) of adult males from their height ( $H$ , in centimeters) and their maximum jump height ( $J$ , in centimeters). We use a model similar to that presented in our previous column<sup>1</sup>, but we now include the effect of  $J$  as  $E(W|H,J) = \beta_H H + \beta_J J + \beta_0 + \varepsilon$ , with  $\beta_H = 0.7$ ,  $\beta_J = -0.08$ ,  $\beta_0 = -46.5$  and normally distributed noise  $\varepsilon$  with zero mean and  $\sigma = 1$  (Table 1). We set  $\beta_J$  negative because we expect a negative correlation between  $W$  and  $J$  when height is held constant (i.e., among men of the same height, lighter men will tend to jump higher). For this example we simulated a sample of size  $n = 40$  with  $H$  and  $J$  normally distributed with means of 165 cm ( $\sigma = 3$ ) and 50 cm ( $\sigma = 12.5$ ), respectively.

Although the statistical theory for MLR seems similar to that for simple linear regression, the interpretation of the results is much more complex. Problems in interpretation arise entirely as a result of the sample correlation<sup>2</sup> among the predictors. We do, in fact, expect a positive correlation between  $H$  and  $J$ —tall men will tend to jump higher than short ones. To illustrate how this correlation can affect the results, we generated values using the model for weight with samples of  $J$  and  $H$  with different amounts of correlation.

Let's look first at the regression coefficients estimated when the predictors are uncorrelated,  $r(H,J) = 0$ , as evidenced by the zero slope in association between  $H$  and  $J$  (Fig. 1a). Here  $r$  is the Pearson correlation coefficient<sup>2</sup>. If we ignore the effect of  $J$  and regress  $W$  on  $H$ , we find  $\hat{W} = 0.71H - 51.7$  ( $R^2 = 0.66$ ) (Table 1 and Fig. 1b). Ignoring  $H$ , we find  $\hat{W} = -0.088J + 69.3$  ( $R^2 = 0.19$ ). If both predictors are fitted in the regression, we obtain  $\hat{W} = 0.71H - 0.088J - 47.3$  ( $R^2 = 0.85$ ). This regression fit is a plane in three dimensions ( $H, J, W$ ) and is not shown in Figure 1. In all three cases, the results of the  $F$ -test for zero slopes show high significance ( $P \leq 0.005$ ).

When the sample correlations of the predictors are exactly zero, the regression slopes ( $b_H$  and  $b_J$ ) for the “one predictor at a time”



**Figure 2** | Results and interpretation of multiple regression changes with the sample correlation of the predictors. Shown are the values of regression coefficient estimates ( $b_H$ ,  $b_J$ ,  $b_0$ ) and  $R^2$  and the significance of the test used to determine whether the coefficient is zero from 250 simulations at each value of predictor sample correlation  $-1 < r(H,J) < 1$  for each scenario where either  $H$  or  $J$  or both  $H$  and  $J$  predictors are fitted in the regression. Thick and thin black curves show the coefficient estimate median and the boundaries of the 10th–90th percentile range, respectively. Histograms show the fraction of estimated  $P$  values in different significance ranges, and correlation intervals are highlighted in red where  $>20\%$  of the  $P$  values are  $>0.01$ . Actual regression coefficients ( $\beta_H$ ,  $\beta_J$ ,  $\beta_0$ ) are marked on vertical axes. The decrease in significance for  $b_J$  when jump height is the only predictor and  $r(H,J)$  is moderate (red arrow) is due to insufficient statistical power ( $b_J$  is close to zero). When predictors are uncorrelated,  $r(H,J) = 0$ ,  $R^2$  of individual regressions sum to  $R^2$  of multiple regression ( $0.66 + 0.19 = 0.85$ ). Panels are organized to correspond to **Table 1**, which shows estimates of a single trial at two different predictor correlations.

regressions and the multiple regression are identical, and the simple regression  $R^2$  sums to multiple regression  $R^2$  ( $0.66 + 0.19 = 0.85$ ; **Fig. 2**). The intercept changes when we add a predictor with a non-zero mean to satisfy the constraint that the least-squares regression line goes through the sample means, which is always true when the regression model includes an intercept.

Balanced factorial experiments show a sample correlation of zero among the predictors when their levels have been fixed. For example, we might fix three heights and three jump heights and select two men representative of each combination, for a total of 18 subjects to be weighed. But if we select the samples and then measure the predictors and response, the predictors are unlikely to have zero correlation.

When we simulate highly correlated predictors  $r(H,J) = 0.9$  (**Fig. 1c**), we find that the regression parameters change depending on whether we use one or both predictors (**Table 1** and **Fig. 1d**). If we consider only the effect of  $H$ , the coefficient  $\beta_H = 0.7$  is inaccurately estimated as  $b_H = 0.44$ . If we include only  $J$ , we estimate  $\beta_J = -0.08$  inaccurately, and even with the wrong sign ( $b_J = 0.097$ ). When we use both predictors, the estimates are quite close to the actual coefficients ( $b_H = 0.63$ ,  $b_J = -0.056$ ).

In fact, as the correlation between predictors  $r(H,J)$  changes, the estimates of the slopes ( $b_H$ ,  $b_J$ ) and intercept ( $b_0$ ) vary greatly when only one predictor is fitted. We show the effects of this variation for all values of predictor correlation (both positive and negative) across 250 trials at each value (**Fig. 2**). We include negative correlation

because although  $J$  and  $H$  are likely to be positively correlated, other scenarios might use negatively correlated predictors (e.g., lung capacity and smoking habits). For example, if we include only  $H$  in the regression and ignore the effect of  $J$ ,  $b_H$  steadily decreases from about 1 to 0.35 as  $r(H,J)$  increases. Why is this? For a given height, larger values of  $J$  (an indicator of fitness) are associated with lower weight. If  $J$  and  $H$  are negatively correlated, as  $J$  increases,  $H$  decreases, and both changes result in a lower value of  $W$ . Conversely, as  $J$  decreases,  $H$  increases, and thus  $W$  increases. If we use only  $H$  as a predictor,  $J$  is lurking in the background, depressing  $W$  at low values of  $H$  and enhancing  $W$  at high levels of  $H$ , so that the effect of  $H$  is overestimated ( $b_H$  increases). The opposite effect occurs when  $J$  and  $H$  are positively correlated. A similar effect occurs for  $b_J$ , which increases in magnitude (becomes more negative) when  $J$  and  $H$  are negatively correlated. **Supplementary Figure 1** shows the effect of correlation when both regression coefficients are positive.

When both predictors are fitted (**Fig. 2**), the regression coefficient estimates ( $b_H$ ,  $b_J$ ,  $b_0$ ) are centered at the actual coefficients ( $\beta_H$ ,  $\beta_J$ ,  $\beta_0$ ) with the correct sign and magnitude regardless of the correlation of the predictors. However, the standard error in the estimates steadily increases as the absolute value of the predictor correlation increases.

Neglecting important predictors has implications not only for  $R^2$ , which is a measure of the predictive power of the regression, but also for interpretation of the regression coefficients. Unconsidered variables that may have a strong effect on the estimated regression coefficients are sometimes called 'lurking variables'. For example, muscle mass might be a lurking variable with a causal effect on both body weight and jump height. The results and interpretation of the regression will also change if other predictors are added.

Given that missing predictors can affect the regression, should we try to include as many predictors as possible? No, for three reasons. First, any correlation among predictors will increase the standard error of the estimated regression coefficients. Second, having more slope parameters in our model will reduce interpretability and cause problems with multiple testing. Third, the model may suffer from overfitting. As the number of predictors approaches the sample size, we begin fitting the model to the noise. As a result, we may seem to have a very good fit to the data but still make poor predictions.

MLR is powerful for incorporating many predictors and for estimating the effects of a predictor on the response in the presence of other covariates. However, the estimated regression coefficients depend on the predictors in the model, and they can be quite variable when the predictors are correlated. Accurate prediction of the response is not an indication that regression slopes reflect the true relationship between the predictors and the response.

*Note: Any Supplementary Information and Source Data files are available in the online version of the paper.*

#### COMPETING FINANCIAL INTERESTS

The authors declare no competing financial interests.

#### Martin Krzywinski & Naomi Altman

- Altman, N. & Krzywinski, M. *Nat. Methods* **12**, 999–1000 (2015).
- Altman, N. & Krzywinski, M. *Nat. Methods* **12**, 899–900 (2015).

Naomi Altman is a Professor of Statistics at The Pennsylvania State University. Martin Krzywinski is a staff scientist at Canada's Michael Smith Genome Sciences Centre.